

# Amazon RDS

COSC 516 – Cloud Databases





# Amazon RDS



**Amazon RDS** (Relational Database Service) is a collection of managed services that simplify operation and scaling of cloud databases.

Supports seven popular database engines:

- Amazon Aurora (MySQL compatible), Amazon Aurora (PostgreSQL compatible)
- MySQL, MariaDB, PostgreSQL, Oracle, SQL Server

Use cases:

- Migrating from on-premise and legacy databases to the cloud
- Reducing maintenance costs by moving to managed databases
- Support growing web and mobile applications

# What is a Managed Database Service?

In a **managed database service** a cloud provider is responsible for:

- Provisioning the hardware and software to run the database
- Ensuring the hardware is always running and working properly
- Installing, configuring, and updating the database server software
- Security, compliance, and software patching
- Provides facilities for performance, scaling, monitoring, backup, and recovery



# Manageability

---

**Monitoring and Metrics:** Use the [RDS Management Console](#) to view operational metrics.

- Additional monitoring and metrics services: [Amazon CloudWatch](#), [Enhanced Monitoring](#), and [Performance Insights](#).

**Event Notifications:** Amazon RDS can notify via email or SMS of events.

# Security

---

**Encryption of Data at Rest:** Data stored in underlying database storage as well as backups and replicas are encrypted. Keys may be managed by [AWS Key Management Service \(KMS\)](#).

**Encryption of Data in Transit:** using SSL/TLS

**Access Control:** using [AWS Identity and Access Management \(IAM\)](#) to control user actions on RDS resources.

**Network Isolation:** using [Amazon Virtual Private Cloud \(VPC\)](#) to isolate DB instances in virtual network and connect using industry-standard encrypted IPsec VPN. May use firewall rules to control access.



# Storage

---

**General Purpose (SSD) Storage:** Consistent baseline of 3 IOPS per provisioned GB and burst up to 3,000 IOPS above the baseline.  
Suitable for a broad range of database workloads.

**Provisioned IOPS (SSD) Storage:** Specify an IOPS rate for the DB instance for predictable and consistent I/O performance.  
Optimized for I/O-intensive transactional (OLTP) database workloads.

# Scalability

---

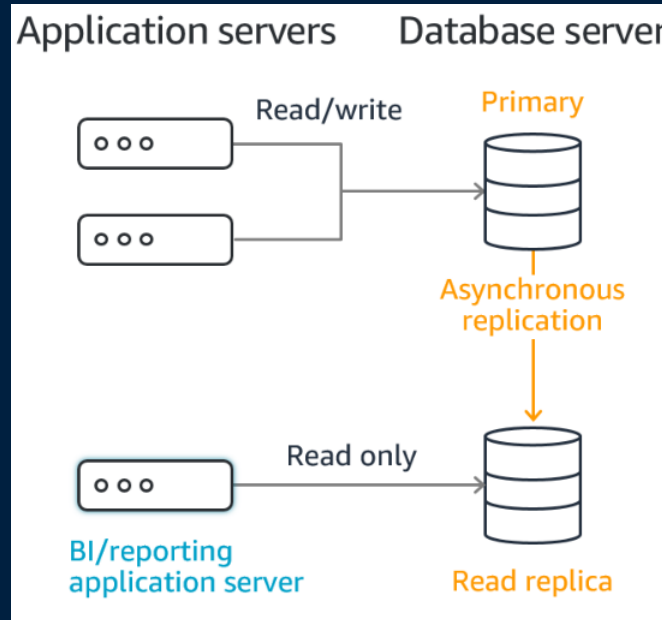
**Compute Scaling:** Adjust compute and memory resources up or down within minutes. Maximum of 32 vCPUs and 244 GiB of RAM.

**Storage Scaling:** Aurora automatically resizes volume up to 64 TB. Other engines require user scaling. Scaling requires no downtime.

**Read Replicas:** DB instance may have read replicas to support heavy read-only traffic and reduce load on primary instance.

# Read Replicas

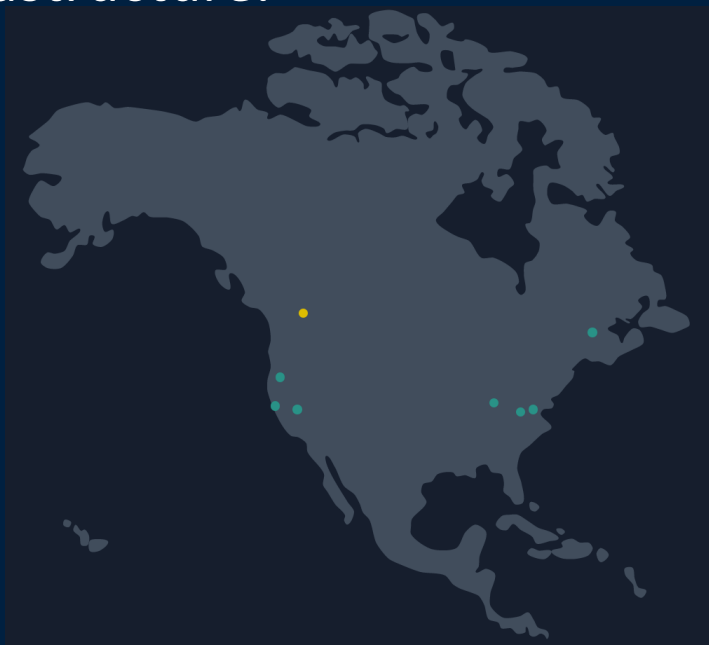
A **read replica** is a database server performing read-only requests for increased performance and availability. The replica is updated from the primary database server using asynchronous replication.





# Availability Zone

An availability zone (AZ) is a distinct location within a region isolated from failures in other AZs. AZ runs on physically distinct, independent infrastructure.



| Region                | AZs    |
|-----------------------|--------|
| US West (Oregon)      | 4      |
| US West (N. Cal)      | 3      |
| US East (N. Virginia) | 6      |
| US East (Ohio)        | 3      |
| GovCloud (US-East)    | 3      |
| GovCloud (US-West)    | 3      |
| Canada (Central)      | 3      |
| Canada (West)         | Coming |

# Amazon Region Map

AWS Cloud has 84 Availability Zones within 26 geographic regions.



Source: Amazon <https://aws.amazon.com/about-aws/global-infrastructure/>

# Replication and High-Availability

---

**Automated Backups:** Backs up database automatically. Backups can be retained up to 35 days.

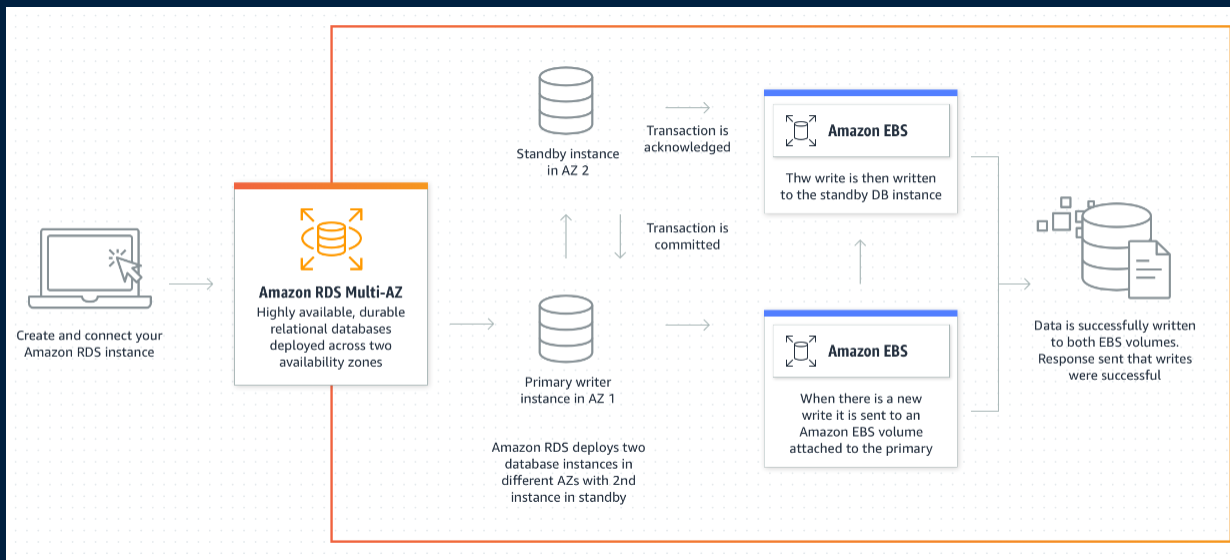
**Database Snapshots:** User-initiated backups stored in Amazon S3. Can use snapshot to create or restore a database instance.

**Multi-AZ Deployments:** Synchronously replicates data to a standby instance in a different Availability Zone (AZ).

**Automatic Host Replacement:** Automatically replaces compute instance hosting database if hardware failure.

# Multi-AZ

A **Multi-AZ** deployment synchronously replicates a DB instance in a different AZ. On failure, automatically fails over to standby instance.



Source: Amazon <https://aws.amazon.com/rds/features/multi-az/>

# MySQL and MariaDB

---

**MySQL** is most popular open source database. RDS supports MySQL Community Edition 5.6, 5.7, and 8.0.

**MariaDB** is an open source database created by the original MySQL developers that is compatible with MySQL.

# Amazon Aurora

---

**Amazon Aurora** is a scalable database designed for the cloud.

Features:

- Protocol compatible with MySQL and PostgreSQL
  - AWS says 5x more throughput than MySQL and 3x more than PostgreSQL
- Storage auto-scaling
- Read replicas share underlying storage (better performance, space utilization)
- Automated multi-region replication
- Continuous backups

**Amazon Aurora Serverless** is on-demand, auto-scaling for Aurora that does not require managing any database instances.

- Useful for variable and unpredictable workloads

# Instances

**DB instance type** defines compute, memory, and storage capabilities.

- E.g. db.t3 is general-purpose burstable instance type (baseline performance but can burst CPU usage as required). db.m6g is general-purpose instance powered by AWS Graviton2 processors.

A **DB instance class** consists of DB instance type and size.

- Sizes: micro, small, medium, large, xlarge
- E.g. db.t3.micro is db.t3 instance type with smallest configuration of resources (1 vCPU, 1 GiB RAM). Available with free tier.

Pricing depends on the DB instance class.



# Cost

---

Cost is based on usage. There is no up-front commitment.

Reserved instances allow for a discount by reserving a instance for a one or three year term.

Cost depends on instance type. Check Amazon RDS Pricing for latest pricing.

The Amazon RDS Free Tier allows 750 hours of Amazon RDS Single-AZ db.t2.micro, db.t3.micro, and db.t4g.micro instances usage running MySQL, MariaDB, PostgreSQL databases each month and 20 GB of General Purpose (SSD) DB storage.

# MySQL Cost – Canada (Central) Region

| Single Deployment | Price Per Hour |
|-------------------|----------------|
| db.t4g.micro      | \$0.018        |
| db.t4g.small      | \$0.035        |
| db.t4g.medium     | \$0.07         |
| db.t4g.large      | \$0.141        |
| db.t4g.xlarge     | \$0.282        |
| db.t4g.2xlarge    | \$0.563        |
| db.t3.micro       | \$0.019        |
| db.t3.small       | \$0.037        |
| db.t3.medium      | \$0.074        |
| db.t3.large       | \$0.148        |
| db.t3.xlarge      | \$0.296        |
| db.t3.2xlarge     | \$0.592        |
| db.m6g.large      | \$0.168        |
| db.m6g.xlarge     | \$0.336        |
| db.m6g.2xlarge    | \$0.672        |

| Multi-AZ Deployment | Price Per Hour |
|---------------------|----------------|
| db.t4g.micro        | \$0.035        |
| db.t4g.small        | \$0.07         |
| db.t4g.medium       | \$0.141        |
| db.t4g.large        | \$0.282        |
| db.t4g.xlarge       | \$0.563        |
| db.t4g.2xlarge      | \$1.126        |
| db.t3.micro         | \$0.037        |
| db.t3.small         | \$0.074        |
| db.t3.medium        | \$0.148        |
| db.t3.large         | \$0.296        |
| db.t3.xlarge        | \$0.592        |
| db.t3.2xlarge       | \$1.184        |
| db.m6g.large        | \$0.336        |
| db.m6g.xlarge       | \$0.672        |
| db.m6g.2xlarge      | \$1.344        |

# MySQL Cost – Canada (Central) Region (2)

---

## Database Storage:

- General purpose: \$0.127 per GB/month
- Provisioned: \$0.138 per GB/month
- Also pay for backup and snapshot storage

## Data Transfer:

- In: Free
- Out: First 10 TB/month - \$0.09 per GB
  - Lower costs if transferring to other Amazon services or regions

# Aurora Cost – Canada (Central) Region

On-Demand instance:

| Standard Instances | Price Per Hour |
|--------------------|----------------|
| db.t4g.medium      | \$0.079        |
| db.t4g.large       | \$0.158        |

Database Storage: General purpose: \$0.11 per GB/month

I/O rate: \$0.22 per 1 million requests

Data Transfer:

- In: Free
- Out: First 10 TB/month - \$0.09 per GB
  - 100 GB free per month per account (applies to all AWS services)

# Conclusion

---

**Amazon RDS** is a managed, relational database service for the cloud.

- It supports seven popular database engines: Amazon Aurora (MySQL and PostgreSQL compatible), MySQL, MariaDB, PostgreSQL, Oracle, SQL Server.

A managed database service handles hardware and software provisioning, maintenance, and availability.

Makes it easy to deploy and scale databases and reduces maintenance time and costs. Amazon RDS is useful for organizations that already have relational databases that want to move to the cloud.

- Database engines function the same as on-premise.
- Costs are a factor that need to be managed.

# Objectives

---

- List the database engines supported by Amazon RDS.
- Compare a managed database service such as Amazon RDS versus deploying a database instance on a cloud server.
- Discuss use cases where a managed database service is valuable.
- Implement security features for Amazon RDS, specifically connecting with SSL.
- Understand when to use general-purpose versus provisioned storage.
- Define read replica and explain when a read replica is useful.
- Define availability zone and region.
- Explain why a database deployed as Multi-AZ has improved durability and availability.
- Explain DB instance type and class and how that relates to resources.
- Compare and contrast Amazon RDS MySQL versus Aurora.
- Understand cloud database costing and how to determine instance pricing.





THE UNIVERSITY OF BRITISH COLUMBIA

