



Graph Data Model



A graph is just a collection of vertices and edges

A set of nodes and the relationships that connect them

Graphs represent entities as nodes and the ways in which those entities relate to the

world as relationships

Label: Person

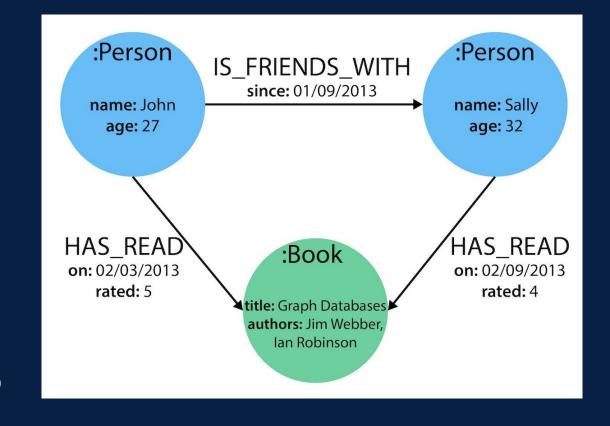
Property: (key, value)

Key: John

Value: 27

Relationships:

IS_FRIENDS_WITH, HAS_READ

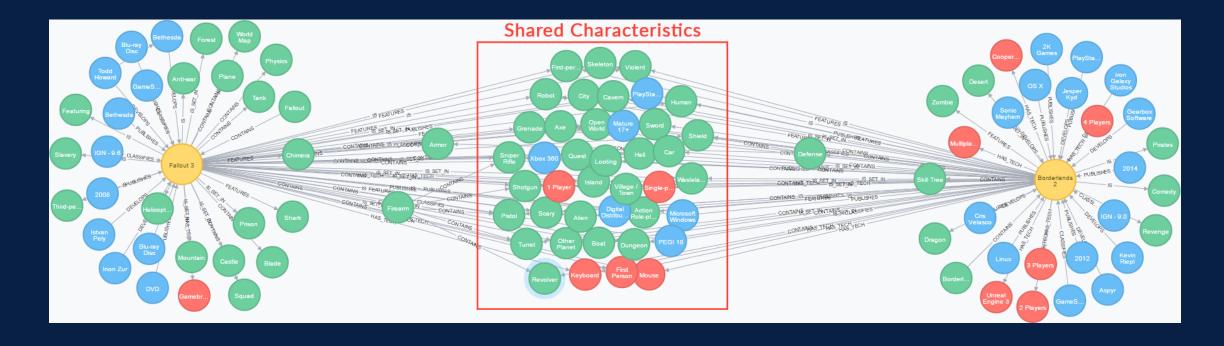






Graph Database Use Cases



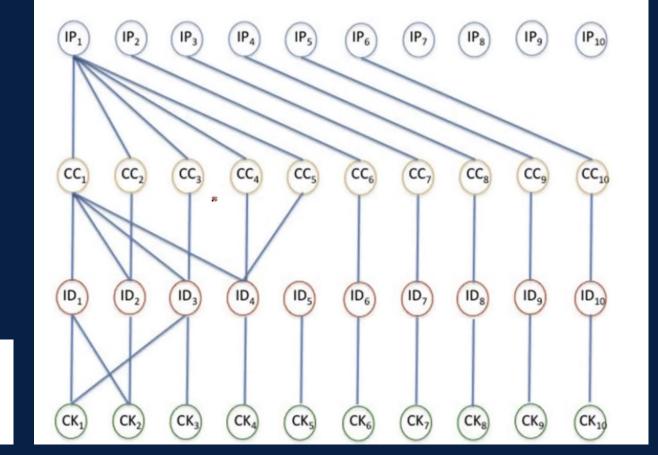


Shared characteristics between video games

Graph Database Use Cases



- Fraud Detection
- Real Time Recommendation Engines
- Master Data Management
- Network and IT Operations
- Identity and Access Management





Walmart

Walmart uses Neo4j to optimize customer experience with personal recommendations

Walmart and eBay adopt graph database

Retail giants turn to database technology to optimise operations, increase insight into customer demand and improve customer service to ultimately drive more sales



Relational Data Model to Graph Data Model



Emp ID	Name	Salary	DNO
112	Sam	1000	1
113	Satish	2000	2

DID	DName		
1	SCOPE		
2	SBST		





Labeled property graph

A labeled property graph has the following characteristics:

- It contains **nodes** and **relationships**
- Nodes contain properties (key-value pairs).
- Nodes can be labeled with one or more labels.
- Relationships are named and directed, and always have a start and end node
- Relationships can also contain properties.



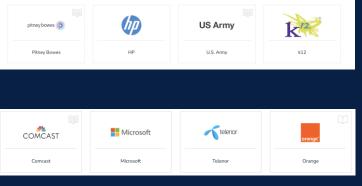


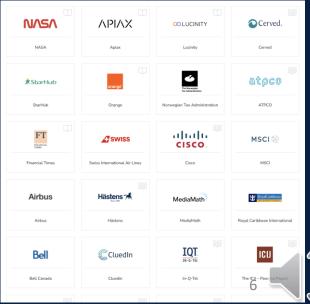


- The development efforts on Neo4j began in 2000, when the co-founders, while on a project building a content
 management system, encountered a problem that couldn't be solved using relational databases so they invented
 the property graph.
- Neo4j is an open-source, NoSQL, native graph database that provides an ACID-compliant transactional backend for your applications that has been publicly available since 2007.
- Neo4j is offered as a managed service via <u>AuraDB</u>. But you can also run Neo4j yourself with either Community Edition or Enterprise Edition. The Enterprise Edition includes all that Community Edition has to offer, plus extra enterprise requirements such as backups, clustering, and failover abilities.

Neo4j Desktop is the new mission control center for developers. It's free with registration, and it includes
a development license for Enterprise Edition.



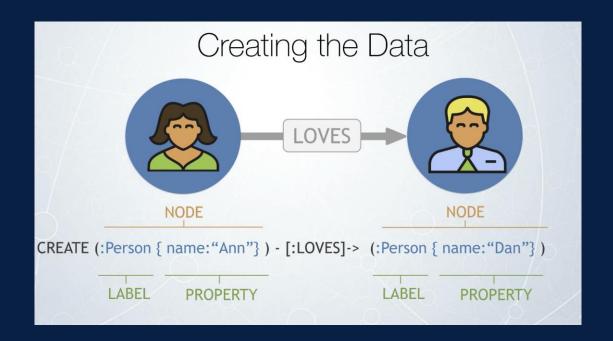


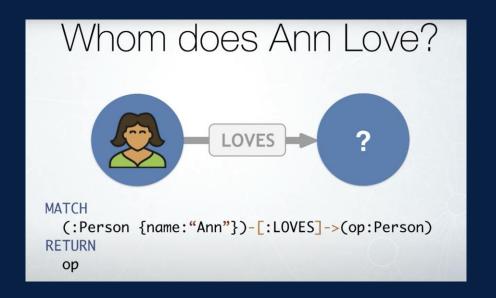






- Cypher is Neo4j's graph query language that allows users to store and retrieve data from the graph database.
- **Cypher is based on patterns**





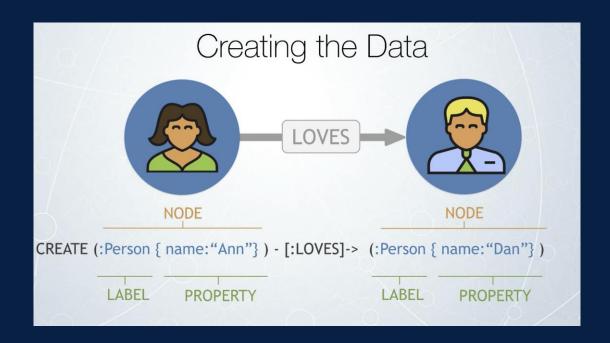
Relationship directions are specified using <, >







- Cypher is Neo4j's graph query language that allows users to store and retrieve data from the graph database.
- **Cypher is based on patterns**



Whom does Ann Love? MATCH (:Person {name: "Ann"})-[:LOVES]->(op:Person) RETURN op It's Dan, of course!

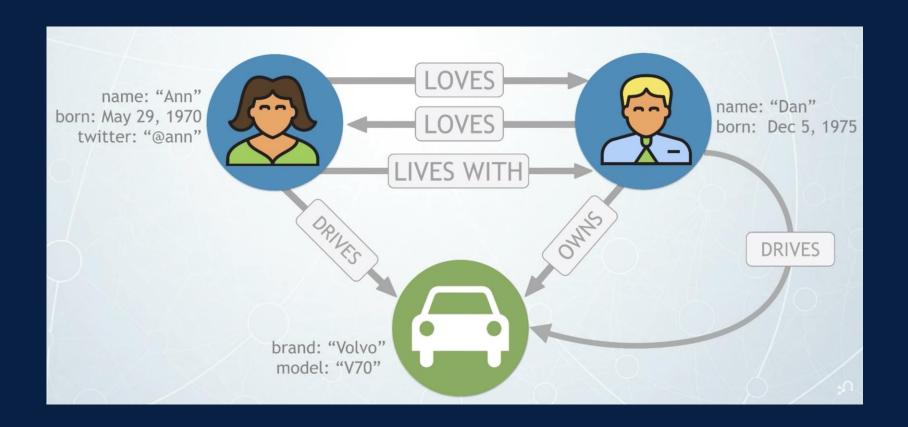
Relationship directions are specified using <, >







Let's say we have now created this graph







```
How do I find Ann's Car?

MATCH
(:Person {name: 'Ann'})-[:DRIVES]->(c:Car)
RETURN
c
```

OR

```
MATCH
(a:Person)-[:DRIVES]->(c:Car)
WHERE
a.name='Ann'
RETURN
C
```

UPDATE Clause

```
MATCH
(a:Person)-[:DRIVES]->(c:Car)
WHERE
a.name='Ann'
SET
c.brand='Volvo',
c.model='V70'
RETURN
c
```

- Add a property to a specific node use SET
- Remove a property for a node use REMOVE
- Update a specific property for a node use REMOVE and SET
- Delete a specific node use MATCH and DELETE
- Delete all nodes match(n) delete(n)
- Delete all nodes with relationships match(n) detach delete n

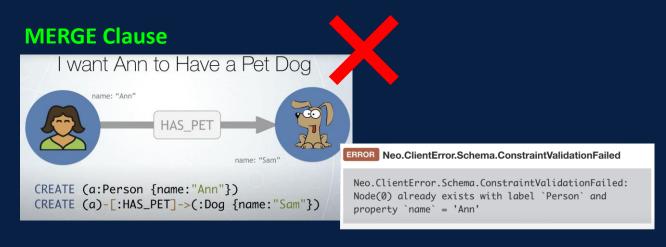


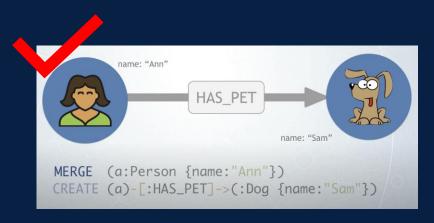


Adding CONSTRAINTS

CREATE CONSTRAINT ON (p:Person)
ASSERT p.name IS UNIQUE

- Unique Constraint
- Property Existence Constraints, eg create constraint on ()-[m:manages]0() assert exists (m.joined)
- Node Key Constraints, eg create constraint on (e:employee) assert (e.empid) is node key





MERGE ON CREATE Clause

- Merge a node and set properties if the node needs to be created
- Example: merge(a:Person {name:"Bill"}) on create set a.born=1996 return a

MERGE ON MATCH Clause

Operators available:
AND, IN, OR (similar to SQL)

- Merging nodes and setting properties on found nodes
- Example: merge(a:Person {name: "Bill"}) on match set a.born=1996 return a



Incoming and Outgoing Relations:

To get all incoming or outgoing relations, use - -> or <- -

Example: match(d:department{name:'COSC'})<- -(e) return d,e</pre>

Matching multiple relationship types:

Use operator in the relationship. Example: match(e)-[m:manages|:worksfor]->(d) return e,m,d

Sort:

Use ORDER BY clause with asc, desc same as SQL. Example: match(e:employee) return e.name order by e.name

Aggregate functions:

All SQL aggregate functions are supported by Cypher – SUM, AVG, COUNT, MIN, MAX Examples: match(e:emp)-[w:worksfor]->(d:dept) return d.name, count(*) match(e:emp) return sum(e.salary)

Other Clauses available:

SKIP - Skips top 'n' records
LIMIT - Returns top 'n' records similar to SQL

(,)

UNION, UNION ALL - Same as SQL, example: match(e:emp{name:'Sam'}) return e.name union
match(e:emp{age:18}) return e.name

STARTS WITH, ENDS WITH, CONTAINS

Source: https://neo4j.com/developer/cypher/intro-cypher/

SQL vs Cypher Query Language



Plsql

```
SELECT DISTINCT c.CompanyName
FROM customers AS c
JOIN orders AS o ON (c.CustomerID = o.CustomerID)
JOIN order_details AS od ON (o.OrderID = od.OrderID)
JOIN products AS p ON (od.ProductID = p.ProductID)
WHERE p.ProductName = 'Chocolade';
```

Cypher

```
MATCH (p:Product {productName: "Chocolade"}) <- [:PRODUCT] - (:Order) <- [:PURCHASED] - (c:Customer)
RETURN distinct c.companyName;
```

```
SELECT e.EmployeeID, count(*) AS Count
FROM Employee AS e
JOIN Order AS o ON (o.EmployeeID = e.EmployeeID)
GROUP BY e.EmployeeID
ORDER BY Count DESC LIMIT 10;
```

```
MATCH (:Order)<-[:SOLD]-(e:Employee)
RETURN e.name, count(*) AS cnt
ORDER BY cnt DESC LIMIT 10
```

```
SELECT p.ProductName
FROM Product AS p
JOIN ProductCategory pc ON (p.CategoryID = pc.CategoryID AND pc.CategoryName = "Dairy Products")
JOIN ProductCategory pc1 ON (p.CategoryID = pc1.CategoryID
JOIN ProductCategory pc2 ON (pc2.ParentID = pc2.CategoryID AND pc2.CategoryName = "Dairy Products")
JOIN ProductCategory pc3 ON (p.CategoryID = pc3.CategoryID
JOIN ProductCategory pc4 ON (pc3.ParentID = pc4.CategoryID)
JOIN ProductCategory pc5 ON (pc4.ParentID = pc5.CategoryID AND pc5.CategoryName = "Dairy Products")
```

```
MATCH (p:Product)-[:CATEGORY]->(1:ProductCategory)-
[:PARENT*0..]-(:ProductCategory {name:"Dairy Products"})
RETURN p.name
```



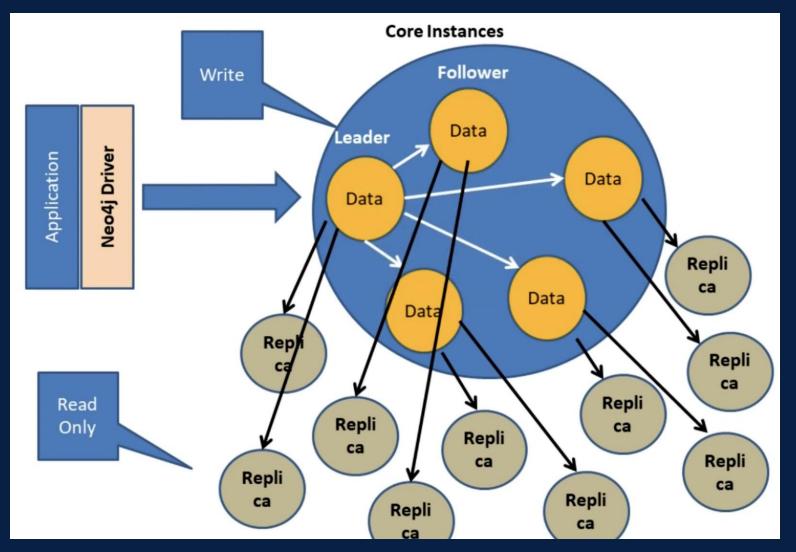


Clusters in Neo4j



Types of Servers in Neo4j Causal Cluster:

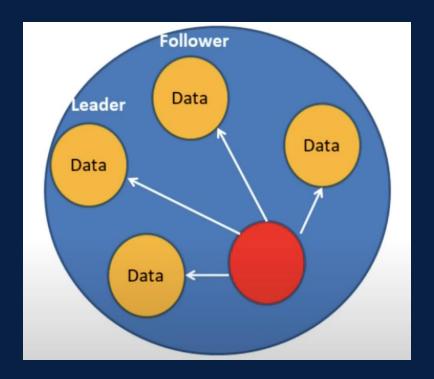
- **Core Servers**
- **Replica Servers**



Discovery Protocol



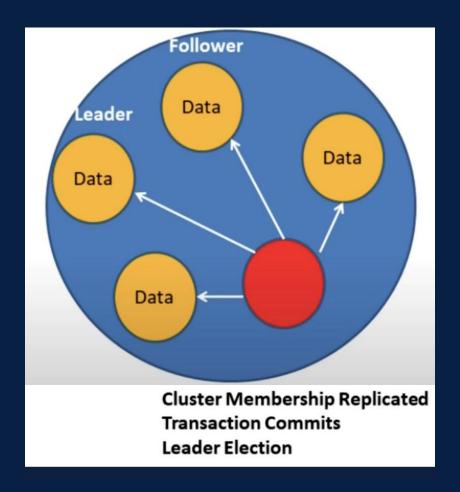
- The discovery protocol takes information from causal_clustering.intital_discovery_ members in neo4j.conf, which lists the IP addresses and ports that form the cluster on start-up.
- When consuming this information, the server will try to handshake with the other listed servers.
- On successful handshake with another server (or servers), the current server will discover the whole current topology.



RAFT Protocol



- Once the Core Server that is performing discovery, has made a connection to one of the existing Core Servers, it then joins the RAFT protocol.
- RAFT is a distributed algorithm for maintaining a consistent log across multiple shared-nothing servers
- RAFT handles cluster membership by making it a normal part of keeping a distributed log in sync.
- Joining a cluster involves the insertion of a cluster membership entry into the RAFT log which is then reliably replicated around the existing cluster.
- Transactions happen via the RAFT protocol. When majority of the instances apply the updates, then the Leader issues commit.
- RAFT protocol also handles leader election.

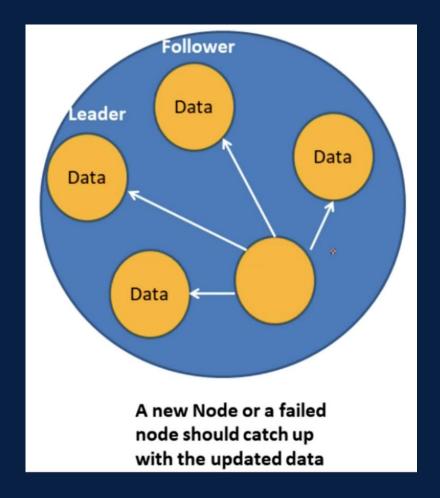




Catch up Protocol



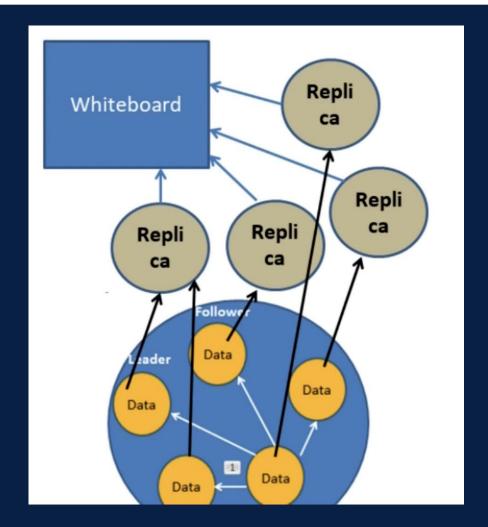
- Keeping the data between the nodes in sync
 - Polling
 - Copying the database store directly



Replica Management



- When a Read Replica performs discovery, once it has made a connection to any of the available Core clusters, it proceeds to add itself into a shared whiteboard.
- This whiteboard provides a view of all live Read Replicas and is used both for routing requests from database drivers that support end-user applications and for monitoring the state of the cluster.
- The whiteboard is kept up to date as Read Replicas join and leave the cluster.

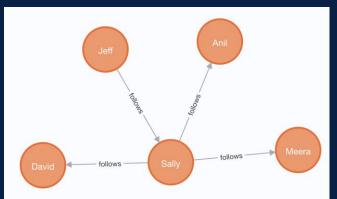




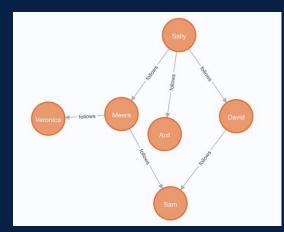
Variable Length Relationships

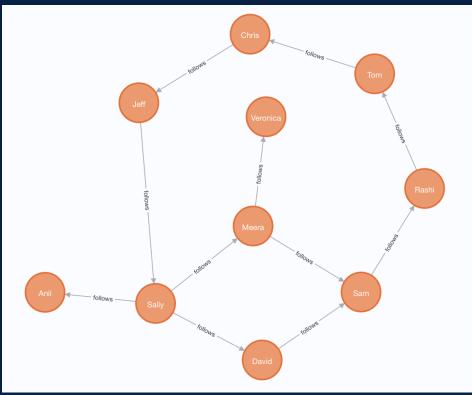


- Nodes that are a variable number of relationship->node hops away, can be found using the following syntax: [:TYPE*minHops..maxHops]->
- minHops and maxHops are optional and default to 1 and infinity respectively.
- When no bounds are given, dots can be omitted.
- The dots may also be omitted when setting only one bound and this implies a fixed length pattern.



lly





Original graph model



Nodes two relationship distance away from Sally

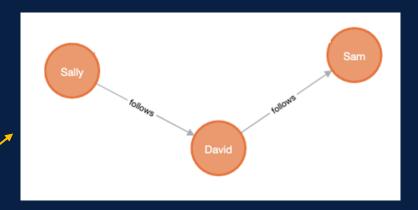
Graph Analytics using Neo4j

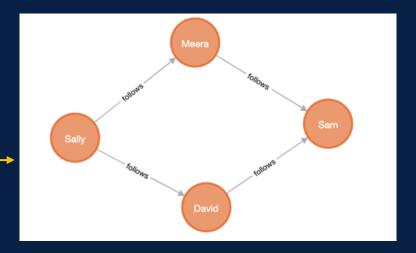


In order to use libraries built for graph analytics, add plugins APOC and GRAPH ALGORITHMS

Shortest Path Algorithm:

- This can be based either on the number of hops or the weighted relationship value.
- What is the shortest path between Sally and Sam with maximum 2 paths away?
 - Hint: Use match with apoc.algo.shortestpath(<condition>) with hop counts
- List all the shortest paths between Sally and Sam:
 - Hint: Use match with apoc.algo.allshortestpath (<condition>)







Graph Analytics using Neo4j

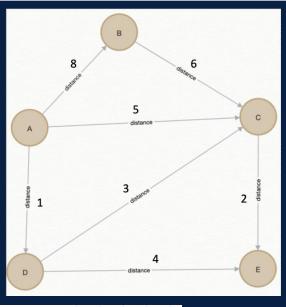


Weighted Graphs

 Determine the weighted shortest path between A and E (ignore directions)

```
match(l:location{name:'A'}),(l1:location{name:'E'})
call algo.shortestPath.stream(l,l1,"cost")
yield nodeId, cost
return algo.getNodeById(nodeId).name, cost
```

"algo.getNodeById(nodeId).name"	"cost"
"A"	0.0
"D"	1.0
"E"	5.0⊳



 Display weighted shortest paths between all pairs of nodes in the graph:

```
call algo.allShortestPaths.stream("cost")
yield sourceNodeId, targetNodeId, distance
where sourceNodeId < targetNodeId
return
algo.getNodeById(sourceNodeId).name,
algo.getNodeById(targetNodeId).name, distance
order by distance</pre>
```



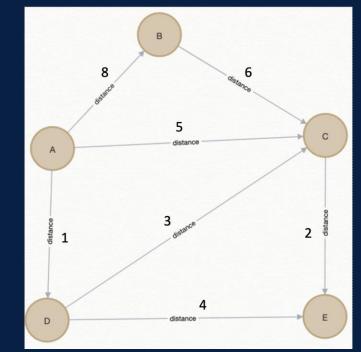


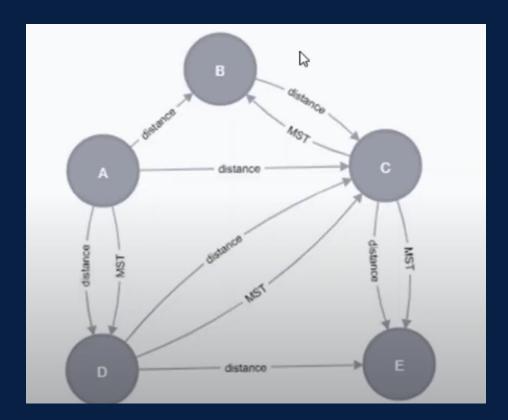
Graph Analytics using Neo4j



Minimum Spanning Tree

- To determine the best route to visit all nodes
- The tree should not contain cycles.
- Determine the MST for the given graph, starting node A.
 - Hint: Use match with algo.spanningTree.minimum("location", "distance", "c ost",id(?))







Cloud Deployments

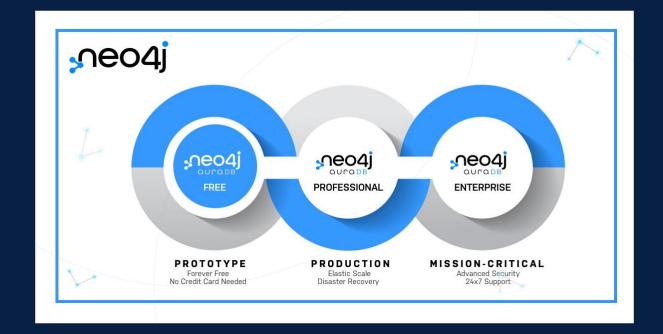




There are different options for deploying Neo4j in the cloud. Neo4j can be deployed in:

- **AWS**
- Google Cloud Platform
- Microsoft Azure

Neo4j AuraDB is a fully managed Neo4j database, hosted in the cloud and requires no installation. https://neo4j.com/cloud/platform/aura-graph-database/

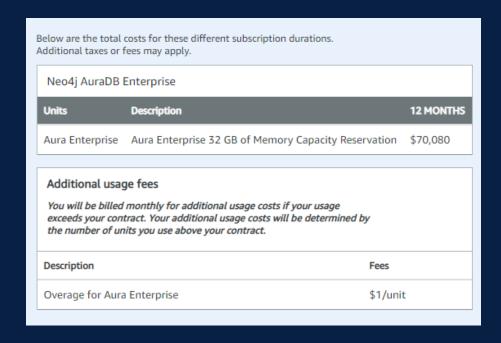


Cloud Deployments

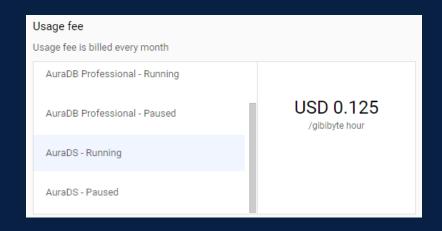




AWS



Google Cloud Platform



Single cluster monthly total: \$476.50 Double cluster monthly total: \$35,536.44

Microsoft Azure

Discovery Bundle	For other configurations, private offers or custom EULAs, contact marketplace-	\$66,000.00/one-time payment	1-year	\$66,000.00
Get it now	sales@neo4j.com	payment		24

Neo4j AuraDB





Neo4j AuraDB – Fully managed cloud database for developers

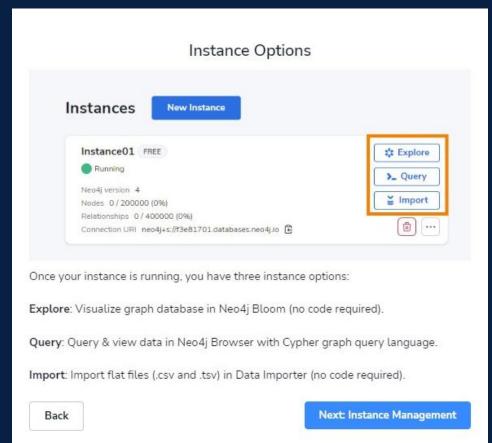
- Fully automated backups, provisioning, upgrades
- Scalable on demand (no interruption)
- Always on secure, reliable, ACID compliant
- Consumption based pricing

Neo4j AuraDS – Graph Data science

Built for data scientists.

Fully managed graph data science engine.
Improves predictions and models at scale to bring projects into production faster. It is deigned to take proof of concept into

production.



Pricing





AuraDB Free

Free

For small development projects, learning, experimentation, and prototyping.

Get Started Free

Feature Highlights:

- Forever free! One fully managed graph database
- No credit card or other payment method required
- Up to 200k nodes and 400k relationships
- More sample datasets to explore
- Bloom data discovery and visualization included
- Access to the largest online graph community

AuraDB Professional

Starting at

\$65/month*

For medium-scale applications in advanced development or production environments.

Sign Up

View Pricing Table

Everything in Free, plus:

- Up to 64 GB memory per database instance
- Unlimited database instances
- Daily backups, 7-day retention
- On-demand snapshots
- Credit card payment or Cloud marketplace billing integration

AuraDB Enterprise



For large-scale, mission-critical applications that require advanced security and 24x7 support.

Contact Us

Everything in Professional, plus:

- Up to 384 GB memory per database instance
- 99.95% guaranteed uptime SLA
- Role-based access control with granular security
- VPC isolation and dedicated infrastructure
- Custom pricing and sales contracts
- ✓ Premium 24x7 support and services
- Hourly backups with 90-day retention
- Support for AWS and GCP marketplace private offers

Cloud - AuraDB

Professional – pay as you go
Enterprise – prepaid contracts with custom pricing
Fast, scalable, graph platform – cloud service

Simple Pricing, Predictable Cost

Simple, transparent pricing

You only pay for provisioned database capacity – no extras for storage, compute, IO, network or backups. With hourly metering, costs are simple to model and predict.

Scale seamlessly, on demand

Easily add or reduce capacity with a few clicks to meet your changing needs, without any service interruption.

Pause on demand, save 80%

Pause your database when not in use and save 80% of running cost. Your data and backups are stored securely. Resume with just a click.

Capacity-based consumption pricing, no extras

Included in your price:

Storage

Backup

⊘IO

O Data Transfer

Pricing





Self – Managed:

- Pricing varies per usage.
- "The best way to run Neo4j onpremises, in your private cloud, or public cloud infrastructure."

Community Edition

For learning Neo4j and smaller projects that do not require high levels of availability, scaling, or professional services and support.

Download for Free

Feature Highlights:

- Open source under GPLv3
- Fully featured, best-in-class native graph database
- Cypher graph query language
- Fully ACID transactions
- Fully compatible with Neo4j Graph Data Science community edition

Enterprise Edition

Enterprise-grade availability and security with scale-up and scale-out options. Run in your private cloud or public cloud infrastructure.

Contact Sales

Feature Highlights:

- Includes all features from community edition
- Unlimited horizontal scaling with replication and sharding
- High availability and advanced manageability
- Available as an operational cluster or stand-alone data science environment

Developer Accessibility





Full suite of developer tools to attract range of users.

Application Development:

- Spring Data Neo4j framework for web applications.
- **GRANDstack** translate GraphQL queries to Cypher.

Easy migration from Neo4j to AuraDB – drag and drop local .dump file.

Graph Apps – visualization and querying with Neo4j, Browser and Bloom, relational database import tools, and monitoring tools.

- Developed by community, partners, enterprises, etc.

Neo4j Officially supported drivers:
NodeJS | Python | JavaScript | Go | .NET | Java

Community drivers:

PHP | Ruby | R | Erlang | Clojure | C/C++ | etc

Neo4j Developer Tools

Neo4j provides sophisticated tools designed to make it easier to develop graph applications.

- Neo4j Desktop is a convenient, free desktop application that allows developers to easily create, work with and manage local Neo4j databases. The free download includes a Neo4j Enterprise Edition license.
- Neo4j Browser lets you interact with your Neo4j database with full CRUD capabilities. You can view results in several formats, including graph visualization mode (for results containing nodes and relationships), tabular and JSON.
- Whether you are just learning about Neo4j or testing out ideas, Neo4j Sandbox is a great way to get your own free, cloud-based instance of Neo4j launched in about 30 seconds, without downloading or installing anything.
- MATCH (c-Customer (email: 'customer@nec4; con' j)' [PURCHASSE) (s-Orders) SETURN co

Developer Tools like Neo4j Desktop, Browser and Sandbox make it easy to learn and develop graph apps.

- Download Neo4j Desktop
- Try Neo4j Sandbox

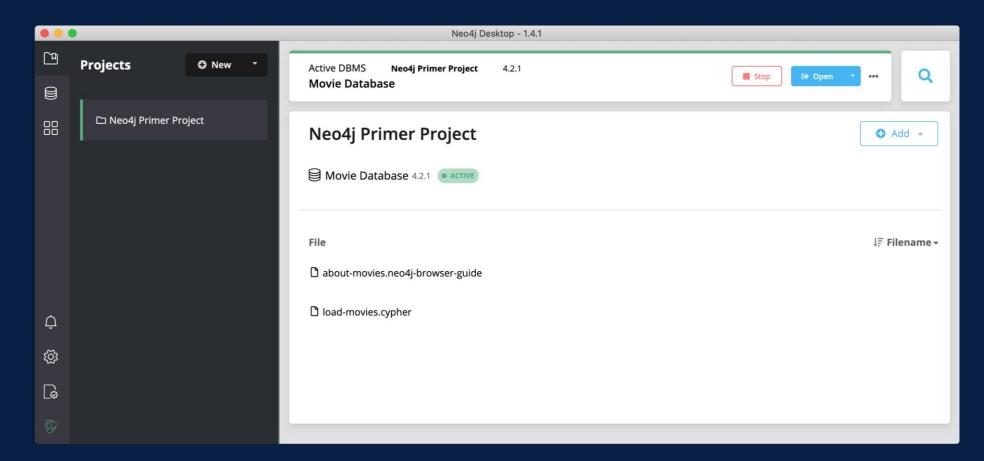
Neo4j Desktop





Neo4j Desktop

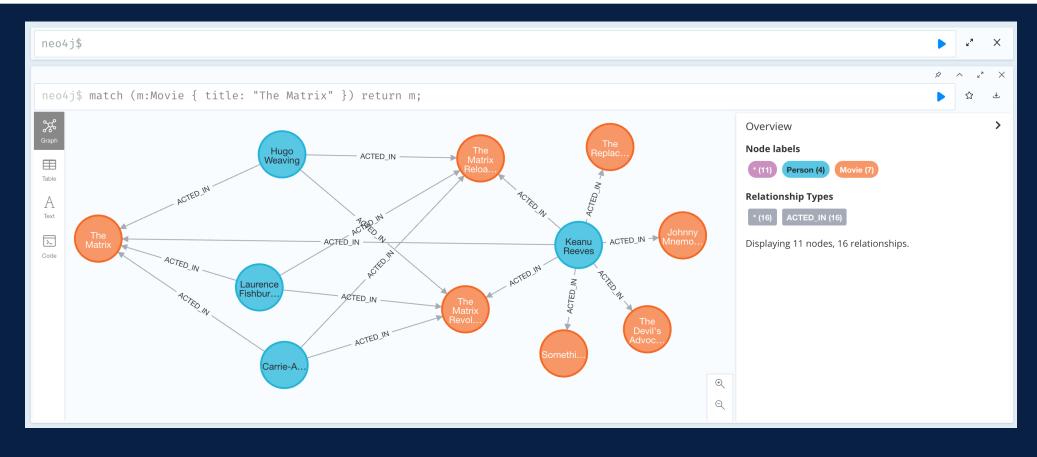
- Developer IDE for Neo4j instances.
- "Similar to Oracle Enterprise Manager."
- Connect and manage local and remote Neo4j servers.
- Comes with free development license of Neo4j Enterprise Edition.



Neo4j Browser







Neo4j Browser

- **Cypher** command shell.
- Interact with graph and visualize data
- Bundled with Neo4j and available in all editions.

Neo4j Bloom

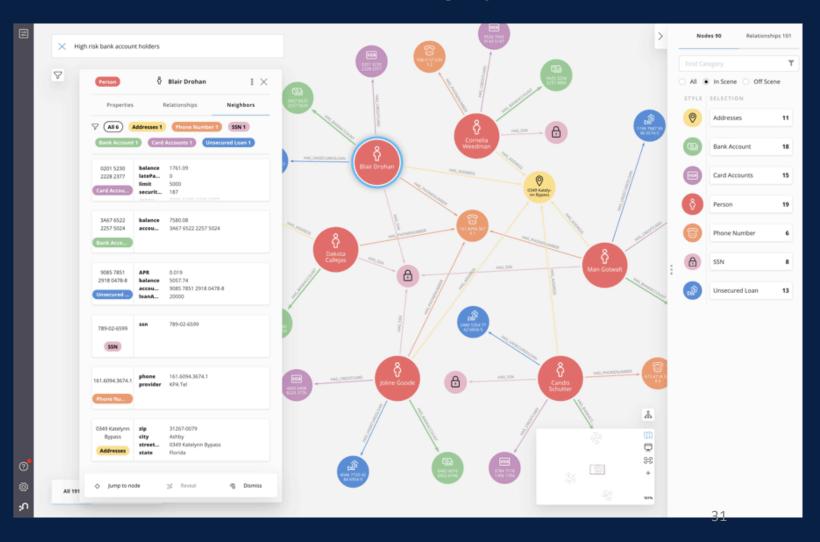


Neo4j Bloom – data exploration tool used to visualize data in the graph.

Allows users to navigate and query data without query language or programming. Users can write patterns to retrieve data.

Accessing Bloom:

- Neo4j Bloom local via Neo4j
 Desktop (local instances).
- Neo4j Bloom server via web browser.
- Neo4j Bloom via sandbox.
- Neo4j Bloom via Neo4j Database (AuraDB).



Neo4j SandBox

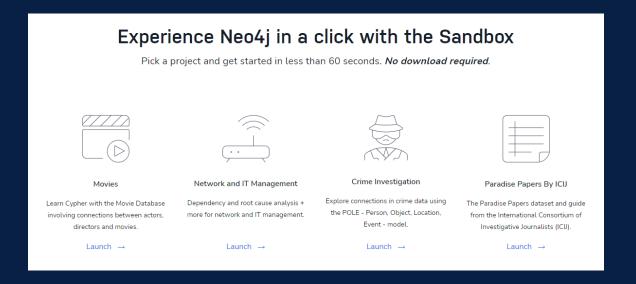


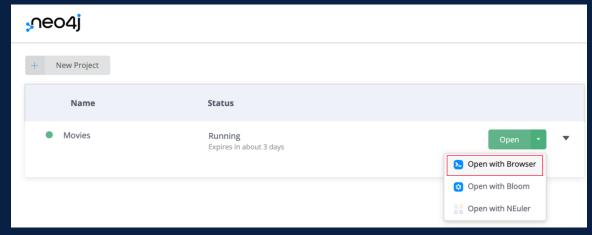


Neo4j Sandbox – free cloud based instance of Neo4j.

- Free. No Download. Fast launch.
- Learn Neo4j, test ideas, gain experience with pre-built data examples.

Launching a project opens Neo4j Browser as a web application and automatically connects to Neo4j instance created by sandbox without and connection URI.

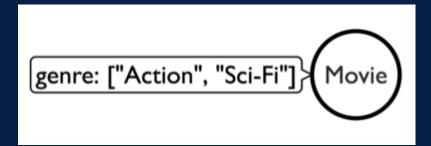




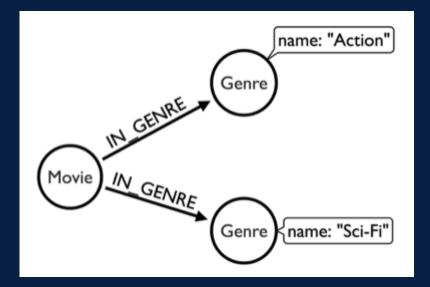
Import Relational Data



Data model design is very important. Neo4j is schema-free meaning your data model can adapt and change. Can easily import CSV files using Cypher to transform contents into graph structure.

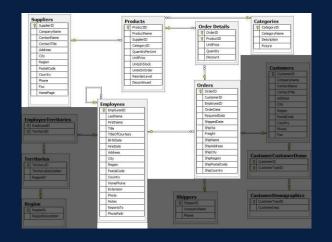


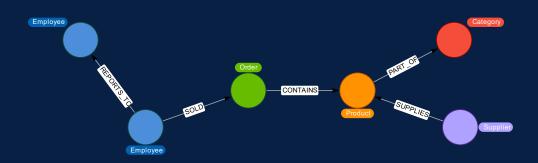
Property on a node vs relationship will determine structure.



Row – node
Table name – label
Joins (fk) – relationship
There are no nulls.

- 1. Create nodes
- 2. Create indexes and constraints on nodes
- 3. Create relationships between nodes
- 4. Querying the graph is ready





Conference



Free educational Neo4j conference.

Keynote speaker.

12 Featured speakers from withing industry/company.
Beginner, Intermediate and advanced agendas.



Join us for this free online graph tech conference for developers and data scientists across the globe. It doesn't matter if you've just heard about graph or you're back to learn best practices – there's something for beginners and experts alike!

November 16 - November 17
See the Agendas Below for Regional Start Times

Keynote Speaker: Nicholas Christakis

Nicholas A. Christakis, MD, PhD, MPH, is the Sterling Professor of Social and Natural Science at Yale University. His work is in the fields of network science, biosocial science, and behavior genetics. He directs the Human Nature Lab and is the Co-Director of the Yale Institute for Network Science. He was elected to the National Academy of Medicine in 2006; the American Association for the Advancement of Science in 2010; and the American Academy of Arts and Sciences in 2017.

Nicholas has also authored these books:

- Connected: The Surprising Power of Our Social Networks and How They Shape Our Lives
- Blueprint: The Evolutionary Origins of a Good Society
- Apollo's Arrow: The Profound and Enduring Impact of Coronavirus on the Way We Live



REGISTER

References



- https://neo4j.com/docs/
- https://neo4j.com/developer/data-import/
- https://www.youtube.com/watch?v=8jNPelugC2s
- https://youtu.be/REVkXVxvMQE

